

Simplified Audio Production in Asynchronous Voice-Based Discussions

Venkatesh Sivaraman
Bexley High School
Columbus, OH, USA
venkats@mit.edu

Dongwook Yoon
Cornell University
Ithaca, NY, USA
dy252@cornell.edu

Piotr Mitros
edX, Inc.
Cambridge, MA, USA
pmitros@edx.org

ABSTRACT

There is no abstract yet.

Author Keywords

Authors' choice; of terms; separated; by semicolons; commas, within terms only; this section is required.

ACM Classification Keywords

H.5.m. Information Interfaces and Presentation (e.g. HCI): Miscellaneous; See <http://acm.org/about/class/1998/> for the full list of ACM classifiers. This section is required.

INTRODUCTION

Asynchronous audio communication (AAC) is rapidly becoming available to mass audiences through social platforms such as WhatsApp, iMessage, and Facebook. While text is still by far the most prevalent mode of communication on the Internet, audio is desirable in many situations because it allows users to deliver more expressive, nuanced messages than text.

AAC holds considerable potential for improving collaboration in education and business as well. For instance, audio messages have been favorably used to disseminate feedback on assignments in online courses [16, 21]. On the other hand, voice communication is much less prevalent in business environments now that voicemail, a paradigm that is generally regarded as “onerous” and “laborious” [9], is on the decline. According to Grudin [11], the primary reason for the failure of voicemail is an inherent imbalance of workload between speaker and listener: message creators can speak more quickly than they can type, but recipients cannot listen as fast as they can read. Perhaps for this reason, email continues to predominate over voice messages in corporate communication despite the diminished individuality and, more critically, the higher risk of miscommunication in textual media [5].

Despite these difficulties, AAC is undoubtedly superior to text for certain applications. For instance, audio messages can be used as annotations on virtual documents, enabling the

user to process comments aurally without visual distraction [35]. Another area of benefit is in online education, where voice communication has been shown to improve student-student and student-instructor engagement as well as a sense of the instructor's social presence [21, 28]. When applying AAC to online education, however, many students may have trouble articulating their ideas vocally. This problem affects students even in physical classrooms and could prevent these learners from participating in online oral discussions. The goal of AAC platforms in such situations, then, is not only to improve efficiency for consumers, but also to compensate for the linearity and immutability of audio on the production side.

Our solution is to provide lightweight, easy-to-use editing tools based on automatic speech recognition (ASR)-generated transcripts. Prior studies have utilized transcription as a proxy for editing audio, but have been limited by the quality of ASR algorithms

We designed and developed a user-friendly audio production tool, which we call SimpleSpeech, and applied it to the scenario of discussion on an online forum. The mental workload experienced by high school students recording voice messages was shown to be significantly decreased with editing functionality, indicating that SimpleSpeech is a valuable enhancement to online audio communication platforms. In addition, some linguistic characteristics of messages created using AAC are discussed with comparison to other forms of computer-mediated communication, leading to new considerations and insights on optimal applications of this technology.

RELATED WORK

Media characteristics of voice comments

Time-aligned transcript for asynchronous voice commenting systems

The linear and serial natures of voice not only preclude skimming and navigation capabilities [11], but also can hamper speech *production* process. Mistakes in a recorded speech is harder to revise than typo in text mainly due to the lack of lightweight editing capability of voice [20]. In addition, producing voice is temporarily linear process which demands the commentator to think and speak simultaneously [20, 36]. Therefore, that the speaker have to keep speaking not to have undesirable long pauses can impose additional cognitive load. Standing on the qualitative implications of these previous works, our study presents quantitative measure about how

Paste the appropriate copyright statement here. ACM now supports three different copyright statements:

- ACM copyright: ACM holds the copyright on the work. This is the historical approach.
- License: The author(s) retain copyright, but ACM receives an exclusive publication license.
- Open Access: The author(s) wish to pay for the work to be open access. The additional fee must be paid to ACM.

This text field is large enough to hold the appropriate release statement assuming it is single spaced.

Every submission will be assigned their own unique DOI string to be included here.

such burdens can be reduced as the voice production system affords the a set of lightweight editing features.

Although voice consumption problem such as the lack of skimming and navigation facilities was well highlighted and tackled by several researchers [33, 26, 19, 35], issues of voice *production* have been documented in only a few studies.

Issues in voice comment production is less documented in the literature compared to the voice access issue in consumption [11] which is originated from the linear nature of a recorded voice. Wimba discussion board study [20]. RichReview CSCW [36]. Grudin[11]. Navigation and skimming.

@Word spotting and pause manipulation [34] envisioned word spotting as a way to provide easy editing and indexing of spoken utterances. [24] envisioned using transcripts as voice indexing and editing aid. Ubiquitous audio [15] and Cedar voice reviewer [1] performs recognition of pauses and speech in a binary fashion to aid visual navigation and editing of audio with and without the presence of speech. It was real-time, but wasn't word level segmentation. Also the editing interface was heavily graphical. Ubiquitous audio [15] and Cedar voice reviewer [1] performs recognition of pauses and speech in a binary fashion. SpeechSkimmer automatically condenses pauses for fast auditory skimming [2]. Pause manipulation using a designated button [3] or tags [23].

@Transcription helps semantic level visual access. For better assess. Previous systems have used time-aligned captions for [29] pitch preserving speed-up and transcription brightness by confidence score for helping consumption. [19] MedSpeak supported editing interface for the transcription errors. Our system allows user to edit both of audio and transcription, but using the cursor position as a clear visual cue for the mode.

@Previous editing systems used time-aligned transcripts as a selection aid. For better production. Previous systems have used time-aligned captions for editing audio [23] or video [8, 3]. Not for production. We first support versatile interactions for live voice production tasks, such as extending pause length, inserting live voice, and lightweight revision of transcription errors.

@text-like editing Yoon et. al's RichReview was another approach but interaction was geared toward pen and touch [35]. Our text-editing like interaction style leverages users familiarity. In radio production scenario [23] did text editing like user interface for editing audio. Tag based p pauses and breaths modulation. Typing ;period; inserted the the pause tags at the point.

@Fixing mistranscript [12] ASR error is detrimental. error fixing strategy. [19] MedSpeak supported editing interface for the transcription errors.

discussion

[3] 's audio removal of 'um's and 'uh's. But it will be challenging in live production scenario where perfect transcription is not guaranteed.

motivation

Mindless jump-cuts results in unnatural pause duration at the boundaries of edited audio tokens.

analysis

0.25 sec cut off. THE TROUBLE WITH "ARTICULATORY" PAUSES* IMPORTANT!! Report our transcription accuracy. Check that it is higher than 84% word error rates, which was evaluated as an OK level [25]. Even errorful transcript is helpful for skimming and understanding [32].

Past approaches to combined audio and text modalities have their origin in voicemail systems, which notoriously favor the producer over the listener. To remedy this, Whittaker et al. [9] attempted to create a more balanced, usable system by converting the audio messages into text using ASR for the recipients. In that study, the combination of text and audio was considerably more appealing to listeners since it allowed them to extract information from messages more efficiently. **Claiming that voice favors producer demotivates our study.**

Other related techniques for audio communication can be found in systems developed for document-based collaboration. Elaborating on the findings of earlier platforms for multimedia annotation [7], an application called RichReview developed by Yoon et al. [35] facilitates discussions on PDF documents via text, audio, or handwritten annotations. Effective use of AAC in these document-based scenarios is especially relevant since it allows commentators to engage in richer discussions over the material without visually distracting from the original content.

In addition to increasing the efficiency of audio consumption, ASR can also be used to facilitate speech editing. Creating effective user interface paradigms that synthesize these two modalities, however, is a challenge that has been addressed mostly in dedicated audio and video editing software.

A similar approach was utilized by Rubin et al. [23], who recently created a text-based editor for "audio stories" containing both speech and music.

In software targeted for non-professional users, Whittaker and Amento [31] developed a "semantic speech editor" in which the sole means of editing was through the transcription. They found that using the transcription as a proxy for audio editing was faster and produced better-quality revisions than audio-only modes, noting that the ability to scan quickly for regions that required editing obviated the need for repeated playback. However, the ASR in their 2004 study suffered from a 28% error rate; today's language models permit considerably greater accuracy, resulting in potentially even easier and faster audio editing.

For instance, a video editor developed by Casares et al. [8] enables the user to edit either the audio track or its transcript side-by-side, with changes in one mode reflected in the other.

DESIGNING SIMPLESPEECH

Building on the capabilities developed in these prior studies, SimpleSpeech is a web-based application for recording and editing short voice messages in a discussion setting. Our design goals were as follows: to utilize live ASR transcription

to assist in speech editing, to enable easy audio manipulation on the word level, and to add more complex features on a tiered basis through modes and quasi-modes. The appearance of the final SimpleSpeech user interface (UI) is shown in Fig. WHAT.

Text-Based Editing

There are a wide variety of approaches to audio editing, ranging from waveform-only interfaces such as Audacity and Adobe Audition to semantic speech editors [31] which show only a transcript. For SimpleSpeech, we decided to adopt an interaction paradigm similar to the latter, allowing the user to edit a textual representation which was time-aligned with the audio. This choice was made for two reasons: (1) in general, users are much more familiar with text than with waveform editing; and (2) representing the audio as text would greatly simplify speech editing on the word level, in contrast to the greater complexity of millisecond-level waveform operations.

Accordingly, the UI for SimpleSpeech devotes the majority of the editing panel to the transcription. Users interact with the text by selecting, editing, and deleting *tokens*, which are colored blue for words and green for pauses and unrecognized sounds. In addition to deleting tokens using the Backspace key, green pause tokens can be inserted or extended using the space bar. Deleting and inserting tokens results in the appropriate modifications automatically applied to the working audio file.

Reinforcing the Connection Between Audio and Text

We chose to include a waveform visualization as part of the UI in order to remind the user that he or she is ultimately manipulating audio, not text. The waveform incorporates several subtle indications of the mapping between its contents and the transcription, including highlighting the audio corresponding to the current selection of tokens and animating deletions and insertions. We found the presence of a waveform to be a helpful visual indicator of the purpose of SimpleSpeech, although it was not functionally useful *per se*. Without the waveform, users' inclination was to disregard the original speech and use the system as a dictation tool.

Another strategy to reinforce the parallelism between the source audio and the transcription is to highlight the words in the transcript as they are spoken during playback. The waveform renders the portion of audio that has already been played in a purple color, which is also used to render the token currently being played back.

Simplification through Modes and Quasi-Modes

Our use of text as a proxy for editing audio rendered it necessary to clearly delineate the capabilities of SimpleSpeech in comparison to a word processor. For instance, direct text input is disallowed in the transcript area to avoid inserting words not present in the original recording. (The inability to move the caret within the tokens visually confirms that the transcript cannot be edited without correspondence to the audio.) However, we found the capability to edit individual words in the transcript to be desirable, especially in the case of ASR errors. The transcription editing functionality is available in a separate mode, accessed by pressing the Return key,

and insulates the editing within single words to avoid undermining the cohesiveness of the tokens. (If the user started typing while one or more tokens were selected, this automatically activated the transcription editing mode as well; some pilot users found this more intuitive than using the Return key.)

During pilot testing the need arose for a fast way to play back and pause the audio; however, the conventional keyboard shortcut for playback, the spacebar, was already in use for the pause insertion feature. We resolved this problem by using Shift+space for playing and pausing. In effect, the playback functionality was encapsulated as a *quasi-mode*, a set of distinct features that are accessed while performing a constant action (in this case, pressing the Shift key) [22]. The modal design helps prevent beginning users from being overwhelmed with possible actions while allowing more advanced users rapid access to the higher-level features.

Pilot Study and Design Improvements

We followed an iterative procedure to develop SimpleSpeech. After building an initial prototype of the application, an informal pilot test was conducted with 5 participants. Each user was given a brief introduction to the software and shown how to use the basic features, then given the scenario of creating an audio response to a written claim on an online forum. (The prompts used in all tests were adapted from the GRE Pool of Issue Topics.) After using the software, users were interviewed to obtain feedback on the prototype, yielding the following modifications:

Tokenization. With the initial tokenization scheme, words and pauses could be deleted in their entirety, but the user could still edit the contents of the tokens by navigating with the arrow keys. However, the pilot study participants were confused by this inline transcript editing behavior; they either tried to insert new unrecorded content by typing or, in the opposite extreme, used the recording insert feature even for minor transcription errors. We endeavored to clarify these delineations in the next iteration by disabling alphanumeric input to the transcript view entirely and changing the transcription editing functionality to a modal interaction.

Pause manipulation. Another important finding in the pilot study was the importance of being able to introduce and adjust pauses between words, not just to remove them. These gaps in the audio help make natural-sounding cuts between audio clips as well as to punctuate claims (e.g., the end of a sentence). The original system only allowed the user to delete pauses, so we added a spacebar action to insert a zero audio signal or fragment of silence from the original audio resource into the rendered message.

Implementation

Our text-based approach requires a reliable transcription as well as time intervals corresponding to each word, both of which are provided by the IBM Watson Developer Cloud speech-to-text transcription service. The web application is written in JavaScript and communicates through a web socket with the IBM server for transcription. Editing is accomplished by maintaining a data model consisting of one or more

user-created audio resources as well as a list of timestamps, each of which links a token in the text area to a time interval within an audio resource. When the user plays back the message, the data model “renders” a complete audio recording by stitching together the audio from each timestamp.

QUALITATIVE EVALUATION

The interaction paradigm of SimpleSpeech was tested in a qualitative assessment to determine (1) the practicability of a lightweight text-based audio editor, (2) the effects of minor transcription errors on audio consumption and production, and (3) the implications of being able to edit audio in an asynchronous online discussion.

Participants were introduced to the functionality of the system, then given two untimed tasks. First, to simulate an asynchronous audio discussion, the test users were asked to listen to an audio comment left by the previous tester and create an original audio response. Next, they received a different, textual prompt and created an audio comment which would be consumed by the next user. In both cases the user was asked to edit his or her recording to be polished and clear. The participants were interviewed at the end of the test; these interviews were transcribed, conversational elements filtered out, and the remaining sentences analyzed via two-step coding (open coding followed by flat coding).

The sample for the study consisted of 9 test subjects (4 male, 5 female; henceforth denoted P_1, P_2, \dots, P_9). All participants were native English speakers. Two individuals, P_2 and P_3 , were professional media editors who provided technical feedback and a comparison to pure audio editing; the remainder were interns and high school students.

Results

The coding process resulted in the following themes identified from the user feedback:

The text-based editing paradigm provides sufficient control to render waveform manipulation unnecessary. Most non-professional users felt SimpleSpeech gave them “plenty of control” over the editing process (P_4, P_5, P_6, P_8). The professional editors did note that most people in their field would not find SimpleSpeech adequate for their needs; but, as P_2 conceded, the intended market users “don’t have to play with the settings which is why they don’t use a professional audio editor.” Most participants characterized the editing experience as being a text-focused one, suggesting that the translation to text was in fact a useful proxy for editing audio. The text modality was described as “more accessible, more doable” than pure waveform editing, which could be “scary for people who don’t do video stuff” (P_3, P_7).

The primary use of lightweight voice editing is to make fine-grained rather than large-scale adjustments. The most commonly-used manipulation during the qualitative study was the removal of disfluencies (P_1, P_2, P_4, P_5, P_7), followed by pause deletion (P_2, P_3, P_5, P_6, P_8). Only P_1 and P_8 edited large chunks of audio by deleting or rerecording, and P_8 reported doing so only to improve the smoothness

of a smaller change in a sentence. Perhaps because SimpleSpeech was presented as a tool to be briefly used to “clean up” recordings, participants focused on removing the “embarrassing” and “awkward” sounds (P_1, P_5).

Transcription is a helpful aid for listening to audio comments despite occasional errors. In many cases, the transcription proved to be an essential element of both the production and the consumption interfaces. To determine the effect of errors in the transcript on listeners, the previous participants’ comments were displayed to users with an unedited, errorful ASR transcript. Despite the occasional errors, users still found the transcript to be helpful in allowing them to “see all the points [the speaker was] making instead of having to remember them” (P_4, P_6). For some users, the transcription caused no problems in comprehension, while others experienced errors that required them to pay more attention to the audio (P_8). On the whole, ASR succeeded in “getting the basic idea across” (P_3) but could not stand alone without the original recording.

The linearity of audio leads to a pressure to organize one’s thoughts during recording. P_4, P_7 , and P_9 described a “psychological sort of ... need to get it all out, and the fact that it won’t necessarily be as organized there.” Another tester, P_5 , had “a tendency to get like a blank slate” in which he “couldn’t think of anything to say.” The elevated mental task load that P_5 describes could be inherent in oral discussion; P_9 noted that “[it] might just be the fact that I was recording,” and that “editing would make it nicer.” Because this phenomenon was present despite the ability to edit, we decided to analyze the task load aspect of using SimpleSpeech in the quantitative study.

Awareness of the recipient and the editability of the audio drive up the quality of contributions. Four users mentioned the formality of their recordings (P_1, P_5, P_7, P_9), which they attributed to “an expectation” to edit, given that “I know that I’ve had that opportunity and someone else would know that I had that opportunity” (P_8). The speakers’ inclination to consider their listeners is exemplified by P_9 , when asked why she was motivated to edit her messages:

Personally I’m editing to express myself a little more in a polished way when I’m writing.... especially if I know someone else is going to review it and be able to respond, I want to make sure I’m as clear as possible and as concise in a way that doesn’t really come across when I’m talking.

Listening to another participant before initiating their own comment may have been a factor in determining the users’ performance, since the exposure “gave ... an understanding of how long of a comment, or what kind of direction people were trying to take the discussion” (P_9). Editing contributed to the increased quality as well: “Since you have the ability to edit things, it feels like you’re talking to somebody who’s prepared a point or a conversational view” (P_5). We chose to explore this phenomenon quantitatively to determine if it was real or simply perceived by the speakers, and to what extent it was affected by the ability to edit.

QUANTITATIVE EVALUATION

For our second, quantitative experiment, we intended to assess the efficacy of SimpleSpeech in particular, and also to measure the usefulness of audio editing tools in general for educational discussions. Participants in the study were given two task parts in random order: recording messages without editing functionality (the No Editing, or NE task) and using SimpleSpeech (the Editing, or E task). Both tasks were presented in a similar UI, except that in the NE part the transcript area was disabled to prevent editing. In each task, users read one of six possible prompt statements, listened to another person's opinion on the issue, then produced an original response. Users took part in two of these "discussion threads" per task. (Before starting the E part, participants were given a standardized tutorial to learn how to edit using SimpleSpeech.)

The quantitative study was conducted at a small public high school, with 16 students (in 11th and 12th grades only) and 14 teachers. This location was ideal for the study because the sample contained a variety of learning and speaking styles; also, the student participants were not all academically oriented, which could have posed a potential bias.

The initial "stimulus" recordings for each of the prompt statements were generated by a group of five initial volunteers. Since the qualitative study had indicated the possibility that prior exposure to other individuals' messages could affect users' perception of formality in the discussion, we divided the stimuli into formal and informal sets. Half the participants of the study (Group A) received only formal recordings, while the other half (Group B) received only informal ones. We hypothesized that the participants in Group A would produce more formal messages due to the stimuli they received.

The criterion used for formality was the F-score, a measure of contextuality introduced by Heylighen and Dewaele in 2002 [14]. The F-score is a purely textual metric based on the frequencies of various parts of speech in a text: nouns, adjectives and prepositions decrease the contextuality and increase the F-score since they are independent of the circumstances around the text, while deictic words such as verbs, adverbs, pronouns, and interjections increase contextuality and decrease the F-score. For our stimulus recordings, the initial participants were asked to plan and edit some of the comments and improvise on the others. After splitting the resulting messages by formality, the average F-score was 53.7 for the Group A messages and 49.4 for the Group B messages, reflecting the greater contextuality of the recordings produced on-the-fly. Group A stimuli also tended to use longer words than those for Group B (4.62 versus 4.38 letters) and tended to be more concise (113 versus 193 words). After obtaining and categorizing these messages, the voices were anonymized by adjusting the pitch randomly.

After each task, the NASA Task Load Index (NASA-TLX) questionnaire was used to quantify the pressure or mental task load of producing a voice message [13]. NASA-TLX is a subjective analytical tool that measures task load along six dimensions: Mental Demand, Physical Demand, Temporal Demand, Performance, Effort, and Frustration. After rat-

ing the level of each aspect of mental workload from 1 (least workload) to 20 (greatest workload), the subject is asked to compare the scales pairwise to produce a weighted TLX value representing the overall pressure during a situation. Participants in the study completed the TLX procedure once after each task to obtain comparisons between the mental workload induced by no-editing and editing situations.

Results

The data collected in the quantitative study fell into three categories, which we will discuss here individually.

Utilization of Editing Features

As in the qualitative study, most participants appreciated and took advantage of the ability to edit their messages. On average, users made about 17 edits to each comment (inserting a new recording, inserting a pause, deleting words, or deleting a pause). Of these changes, the vast majority were subtractive: 7 word deletions and 6.3 pause deletions per message. This was again consistent with the findings of the earlier study, which had shown an inclination to remove disfluencies and "awkward" hesitations from the recordings.

The editing showed significant differences in how time was spent in the resultant recordings: (give editing statistics)

Ideally, participants in the E task would have edited both the transcription and the voice to be free of errors, but due to time constraints on participation we discouraged the users from correcting transcript errors (which turned out to be more numerous than expected, especially because of the conversational style). Incorrect transcriptions were problematic for editing in general: Since the associated timestamps were also incorrect, the edits on that segment of audio could produce undesirable results.

A few participants found the mechanism to edit the transcript (pressing Return) somewhat counterintuitive. When it was presented in the tutorial, at first these participants tried to use the Delete key on the token they wanted to edit, resulting in the permanent deletion of the word and its corresponding audio. The emphasis in the tutorial that the Delete key deleted the audio permanently did help other participants avoid making this mistake.

Another misconception we observed in a few participants was a tendency to treat SimpleSpeech as a dictation tool. These users paused for long periods of time during recording sessions and neglected to play back the messages during editing. Furthermore, their inclination after stopping a recording session was to go back and correct transcription errors so that the visual representation made sense.

Task Load

Since the NASA-TLX scale is subjective, it does introduce variability between participants due to the differences between their perceived skill at the task [13]. For instance, one participant could rate the recording task at a 3 out of 20, while another could rate the very same task at a 15. Therefore, the strongest comparisons of task load were made in the within-subject dimension, which was the ability or inability to edit.

Students ($N = 16$)			
Task	<i>E</i>	<i>NE</i>	
Mental Demand	9.6	11.1	
Physical Demand	3.7	2.6	
Temporal Demand	7.8	10.5	*
Performance	8.3	10.0	*
Effort	9.1	11.6	*
Frustration	7.8	8.9	
Total (weighted)	8.7	10.8	*

Table 1. The mental work load ratings reported by students, along with the total task load index values. (* – $p < 0.05$)

Overall, the students reported significantly *lower* levels of mental task load or pressure during the E task than the NE task (mean 8.7 compared to 10.8, $p < 0.02$). The values for the individual components of the TLX, shown in Table 1, yielded the following significant contributory dimensions on the TLX questionnaire:

- *Temporal demand.* Students rated the temporal demand at 7.8 for the E task, compared to the NE rating of 10.2. As described by the TLX form, temporal demand refers to “time pressure due to the rate or pace at which the tasks or task elements occurred” [13]. Students verbally described the increase in time demand reported on the TLX in terms of having to think of words quickly, with the knowledge that every second not filled with speech would be an embarrassing silence.
- *Performance.* Students felt significantly more concern about the quality of their messages in the NE task, rating it at 10.3 compared to 8.0 for the E task. Just as the participants in the prior qualitative study had articulated a desire to make their messages better for the sake of their listeners, the students also evidently wanted to improve their recordings in the NE task. The inability to do so resulted in elevated task load due to performance, while for the E task the stress was lower because they were afforded the chance to correct their mistakes. However, it is worth noting that even despite the capability to edit, the student participants still rated Performance close to the middle of the scale, perhaps representing self-consciousness or comparisons with the stimulus recordings.
- *Effort.* Similarly to performance, students reported having to work harder in the NE task to complete it to their desired level (rated 12.1 compared to 9.0 in the E task). This increased effort could correspond to the additional mental activity which had to be expended in order to generate speech fluently and without excessive hesitation.

The teachers also reported slightly lower average workload levels in the E task, as shown in Table 2, but this difference was not significant. In fact, 7 of the 13 participating teachers actually rated the E task as requiring a higher workload than the NE task. This subset of the teachers, 5 of whom were in Group A, reported an average task load greater in the E task than the NE task for *all* dimensions, especially Mental

Teachers ($N = 13$)			
Task	<i>E</i>	<i>NE</i>	
Mental Demand	11.5	11.2	
Physical Demand	4.0	3.2	
Temporal Demand	7.8	9.9	
Performance	8.2	9.5	
Effort	10.0	10.4	
Frustration	8.4	9.7	
Total (weighted)	9.5	10.6	
$p > 0.05$			

Table 2. The mental work load ratings reported by teachers, along with the total task load index values.

Demand, Performance, and Frustration. The reason for this rating, these teachers explained, was that the availability of the editing tools caused them to feel more worried about their performance. Editing in turn required them to expend more effort to preserve the existing fluidity of their messages.

Interestingly, this particular group of teachers produced more formal messages than the other teachers (mean F-score 56.6 compared to 53.1), with longer words (4.58 compared to 4.36 letters), and fewer disfluencies (1.3 compared to 2.1 per 100 words). Upon further inspection, the student participant group also contained members who rated the E task as more demanding than the NE task, though fewer in number (4 out of 16); these students also produced much more formal messages than their peers (57.9 compared to 53.5). These participants could have had more experience speaking extemporaneously or felt less inclined to speak conversationally, ultimately leading to SimpleSpeech not being as useful to them.

Overall, the fact that the differences in perception of workload varied so much among teachers indicates that they were not as heavily affected by the ability to edit as the students, who clearly appreciated the security that SimpleSpeech offered.

Formality

Contrary to the hypothesis that prior exposure to audio messages would affect the formality or linguistic traits of new messages, the F-scores of the participants’ output was unrelated to the group they were in, as shown in Figure 3. As in the workload measurements, the average F-score for students was higher for Group A (55.83) than for Group B (53.42), which is a considerable difference in terms of the F-score’s scale but not statistically significant. The F-scores for teachers were almost distinguishable, with a difference of only 0.38. In other words, the formality of the recordings was not affected by the stimulus message or even whether the participant was a teacher or a student. Considering that the F-score measures contextuality between the speaker and the audience, and that its value was not affected by the context given before the tasks, the principal source of variation in F-score must have been personal preference in the medium and the scenario of an online forum discussion.

	Students		Teachers	
	A	B	A	B
Group				
Formality (F-score)	55.83	53.42	54.74	55.12
Word Length	4.40	4.44	4.50	4.46
Disfluencies (per 100 words)	1.59	2.38	1.27	1.95
Word count	100.66	140.47	155.38	127.00
Speaking rate	130.39	114.75	136.58	139.01

Table 3. Various metrics representing the formality of the audio messages produced by each participant group.

FORMALITY COMPARISON

Given that the F-scores of SimpleSpeech messages were roughly normally distributed and not heavily affected by the experimental conditions, the average F-score of 54.8 is likely to be characteristic of general AAC discussions under similar conditions. Contextuality in the online voice-based forum scenario could be highly indicative of AAC’s potential applications, and to our knowledge this trait has not been studied extensively. The closest related studies have pertained to other forms of computer mediated communication (CMC), especially textual ones such as SMS, email, or Facebook posts. For example, Kiesler, Siegel, and McGuire [17] found more equalized group participation and more uninhibited expression of opinions in synchronous text-based CMC than in face-to-face discussions. Asynchronous CMC, similar to a discussion board, induces more prosocial behavior and, in fact, more informal communication styles over time than face-to-face [30]. On the other hand, formality and politeness in emails has been shown to increase as the social distance, status gap, and importance of a request increase [6].

How AAC fits into the complex hierarchy of social dynamics on various platforms is still unknown, so we conducted a comparison of the voice messages composed during this study with corpora of different media. The SimpleSpeech text, the focal point of the comparison, contained – words from – messages. For written documents, we used several sections of the well-known Brown corpus to compile general categories of text: nonfiction, fiction, and technical writing (consisting of government documents, scientific articles, and news) [18]. We obtained chatroom text from the `nps_chat` corpus, face-to-face conversation data from the `webtext` corpus, and telephone data from the `switchboard` corpus, all available as part of the Natural Language Toolkit (NLTK) [4]. Finally, we also analyzed email communication in non-spam messages from the Enron corpus [27], as well as a corpus of Twitter posts [10].

Results

The results of this comparison, shown in Figure 1, illustrate the middle-ground that AAC takes relative to oral and written media. The least formal and most contextual corpora were those based on oral communication (with the notable exception of web chat messages), while the most formal and least context-dependent were the written texts, including email and Twitter posts. We will note three additional explanations for




Figure 1. The formality of corpora in different genre and media. The messages produced using SimpleSpeech during the quantitative study are intended to reflect general AAC discussion characteristics.

the formality of each medium based on the ordering of the corpora:

Speaker-audience relationship. Since the F-score is inversely related to contextuality, it is reasonable that the chat and telephone corpora had the lowest F-scores because the participants knew each other and were conversing on a one-to-one basis. On the other hand, the written forms of communication (with the exception of email) were more formal because the audience was defined more loosely and not necessarily acquainted with the speaker. AAC using SimpleSpeech was more closely related to the latter condition (as an online forum discussion), which probably contributed to its greater formality compared to the other spoken corpora.

Immediacy of communication. The tendency to speak or write more contextually when the recipient replies immediately explains why the online chat text, though written, was more contextual and less formal than the oral corpora. It also justifies the fact that the email corpus was more formal than all of the other direct communication media. Again, AAC falls toward the more formal end of this spectrum because there is little temporal proximity between the speaker and the audience.

Tendency toward verbosity. Media that pressured the creator to be brief or precise were more formal and less contextual. For instance, writing technical documents requires the preferential use of nouns over pronouns to maximize clarity. Twit-

ter messages are, of course, limited to 140 characters, leading to a greater concentration of meaning that favors less contextual words. For AAC, therefore, the ability to edit could influence the contextuality if discussion members were pressured to trim down their recordings. For our study, however, the participants were not affected by verbosity; though non-edited recordings had on average 10% more words than edited ones, these edits were more concentrated on removing disfluencies than improving concision.

ACKNOWLEDGMENTS

Sample text: We thank all the volunteers, and all publications support and staff, who wrote and provided helpful comments on previous versions of this document. Authors 1, 2, and 3 gratefully acknowledge the grant from NSF (#1234–2012–ABC). *This whole paragraph is just an example.*

REFERENCES

1. Stephen Ades and Daniel C Swinehart. 1986. *Voice annotation and editing in a workstation environment*. XEROX Corporation, Palo Alto Research Center.
2. Barry Arons. 1993. SpeechSkimmer: Interactively Skimming Recorded Speech. In *Proceedings of the 6th Annual ACM Symposium on User Interface Software and Technology (UIST '93)*. ACM, New York, NY, USA, 187–196. DOI : <http://dx.doi.org/10.1145/168642.168661>
3. Floraine Berthouzoz, Wilmot Li, and Maneesh Agrawala. 2012. Tools for Placing Cuts and Transitions in Interview Video. *ACM Trans. Graph.* 31, 4, Article 67 (July 2012), 8 pages. DOI : <http://dx.doi.org/10.1145/2185520.2185563>
4. Steven Bird, Ewan Klein, and Edward Loper. 2009. *Natural Language Processing with Python*. O'Reilly Media.
5. Kristin Byron. 2008. Carrying Too Heavy a Load? The Communication and Miscommunication of Emotion by Email. *Academy of Management Review* 33, 2 (2008).
6. Thomas Cho. 2010. Linguistic Features of Electronic Mail in the Workplace: A Comparison with Memoranda. *Language@Internet* 7, 3 (2010).
7. David Barger et al. 2002a. Asynchronous Collaboration Around Multimedia Applied to On-Demand Education. *Journal of Management Information Systems* 18, 4 (2002).
8. Juan Casares et al. 2002b. Simplifying Video Editing Using Metadata. In *DIS '02 Proceedings*. London, 157–166.
9. Steve Whittaker et al. 2002c. SCANMail: a voicemail interface that makes speech browsable, readable and searchable. *CHI Letters* 4, 1 (2002).
10. Alec Go, Richa Bhayani, and Lei Huang. 2009. *Twitter Sentiment Classification using Distant Supervision*. Technical Report. Stanford.
11. Jonathan Grudin. 1988. Why CSCW applications fail: Problems in the design and evaluation of organizational interfaces. *ACM Conference on Computer-Supported Cooperative Work* (1988).
12. C Halverson, D Horn, C Karat, and John Karat. 1999. The beauty of errors: Patterns of error correction in desktop speech systems. In *Proceedings of INTERACT99*. 133–140.
13. Sandra G. Hart and Lowell E. Staveland. 1988. Development of NASA-TLX (Task Load Index): Results of empirical and theoretical research. In *Human Mental Workload*, P. A. Hancock and N. Meshkati (Eds.). North Holland Press, Amsterdam.
14. Francis Heylighen and Jean-Marc Dewaele. 2002. Variation in the contextuality of language: an empirical measure. *Foundations of Science* 7 (2002), 239–340.
15. Debby Hindus and Chris Schmandt. 1992. Ubiquitous Audio: Capturing Spontaneous Collaboration. In *Proceedings of the 1992 ACM Conference on Computer-supported Cooperative Work (CSCW '92)*. ACM, New York, NY, USA, 210–217. DOI : <http://dx.doi.org/10.1145/143457.143481>
16. Philip Ice, Reagan Curtis, Perry Phillips, and John Wells. 2007. Using asynchronous audio feedback to enhance teaching presence and students' sense of community. *Journal of Asynchronous Learning Networks* 11, 2 (2007).
17. Sara Kiesler, Jane Siegel, and Timothy W. McGuire. 1984. Social Psychological Aspects of Computer-Mediated Communication. *Amer. Psychologist* 39, 10 (1984), 1123–1134.
18. Henry Kučera and W. Nelson Francis. 1967. *Computational Analysis of Present-Day American English*. Brown University Press, Providence.
19. Jennifer Lai and John Vergo. 1997. MedSpeak: Report Creation with Continuous Speech Recognition. In *Proceedings of the ACM SIGCHI Conference on Human Factors in Computing Systems (CHI '97)*. ACM, New York, NY, USA, 431–438. DOI : <http://dx.doi.org/10.1145/258549.258829>
20. Philip Marriott. 2002. Voice vs text-based discussion forums: An implementation of Wimba Voice Boards. In *World Conference on E-Learning in Corporate, Government, Healthcare, and Higher Education*, Vol. 2002. 640–646.
21. Jody Oomen-Early, Mary Bold, Kristin L. Wiginton, Tara L. Gallien, and Nancy Anderson. 2008. Using asynchronous audio communication (AAC) in the online classroom: a comparative study. *Journal of Online Learning and Teaching* 4, 3 (2008).
22. Jef Raskin. 2000. *Humane Interface, The: New Directions for Designing Interactive Systems*. Addison-Wesley.

23. Steve Rubin, Floraine Berthouzoz, Gautham J. Mysore, Wilmot Li, and Maneesh Agrawala. 2013. Content-Based Tools for Editing Audio Stories. In *UIST '13*. 113–122.
24. Christopher Schmandt. 1981. The Intelligent Ear: A Graphical Interface to Digital Audio. In *Proceedings, IEEE International Conference on Cybernetics and Society, IEEE*.
25. Litza A Stark, Steve Whittaker, and Julia Hirschberg. 2000. ASR satisficing: the effects of ASR accuracy on speech retrieval.. In *INTERSPEECH*. Citeseer, 1069–1072.
26. Lisa Stifelman, Barry Arons, and Chris Schmandt. 2001. The Audio Notebook: Paper and Pen Interaction with Structured Speech. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '01)*. ACM, New York, NY, USA, 182–189. DOI : <http://dx.doi.org/10.1145/365024.365096>
27. Will Styler. 2011. *The EnronSent Corpus*. Technical Report 01-2011. University of Colorado at Boulder Institute of Cognitive Science, Boulder, CO.
28. Chi-Hsiung Tu and Marina McIsaac. 2002. The Relationship of Social Presence and Interaction in Online Classes. *Amer. Journal of Distance Education* 16, 3 (2002).
29. Sunil Vemuri, Philip DeCamp, Walter Bender, and Chris Schmandt. 2004. Improving Speech Playback Using Time-compression and Speech Recognition. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '04)*. ACM, New York, NY, USA, 295–302. DOI : <http://dx.doi.org/10.1145/985692.985730>
30. Joseph B. Walther. 1995. Relational Aspects of Computer-Mediated Communication: Experimental Observations over Time. *Organization Science* 6, 2 (1995), 186–203.
31. Steve Whittaker and Brian Amento. 2004. Semantic Speech Editing. *CHI Letters* (2004), 527–534.
32. Steve Whittaker, Julia Hirschberg, John Choi, Don Hindle, Fernando Pereira, and Amit Singhal. 1999. SCAN: Designing and Evaluating User Interfaces to Support Retrieval from Speech Archives. In *Proceedings of the 22Nd Annual International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '99)*. ACM, New York, NY, USA, 26–33. DOI : <http://dx.doi.org/10.1145/312624.312639>
33. Steve Whittaker, Patrick Hyland, and Myrtle Wiley. 1994. FILOCHAT: Handwritten Notes Provide Access to Recorded Conversations. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '94)*. ACM, New York, NY, USA, 271–277. DOI : <http://dx.doi.org/10.1145/191666.191763>
34. Lynn Wilcox, Ian Smith, and Marcia Bush. 1992. Wordspotting for Voice Editing and Audio Indexing. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '92)*. ACM, New York, NY, USA, 655–656. DOI : <http://dx.doi.org/10.1145/142750.150715>
35. Dongwook Yoon, Nicholas Chen, François Guimbretière, and Abigail Sellen. 2014. RichReview: blending ink, speech, and gesture to support collaborative document review. In *Proceedings of the 27th annual ACM symposium on User interface software and technology*. ACM, 481–490.
36. Dongwook Yoon, Nicholas Chen, Bernie Randles, Amy Cheatle, Corinna E. Loeckenhoff, Steven J. Jackson, Abigail Sellen, and François Guimbretière. 2016. Deployment of a Collaborative Multi-Modal Annotation System for Instructor Feedback and Peer Discussion. In *Proceedings of the 19th ACM Conference on Computer Supported Cooperative Work & Social Computing*. ACM.